

HONEYWELL DCS HISTORIAN ANALYSIS

Advanced Process Exploratory Data Analysis (EDA) Report

Prepared by: Data Science & Process Analytics Team

Date: June 2026

Target Tag: 03TIC_1023.PV (Column Overhead Temperature)

Safety Threshold: PVHI $\geq$ 21.0 °C

Data Span: Jan 2022 to Jan 2026 (4 Widescreen Years)

Data Volume: 2,019,221 continuous records (minutes)

Executive Summary

This diagnostic report provides a comprehensive evaluation of the historical process values associated with the overhead distillation column temperature alarm (tag 03TIC_1023.PV). Over the 4-year continuous operational period, the column recorded 1,561 safety violations where the temperature met or exceeded the PVHI threshold of 21.0°C. By applying advanced exploratory data analysis methods—specifically consecutive null audits, triple correlation matrices (Pearson, Spearman, and Distance correlation), alarm duration buckets, and pre/post-alarm ramping dynamics—we uncovered critical process characteristics. These insights directly inform the design of our multi-horizon predictive alarm models and identify clear process bottlenecks.

Key Findings:


- **Data Quality & Imputation:** Missing data rates are extremely low (<0.15%). Since all tags fell well below the 60% missingness threshold, a 5-minute capped forward-fill was successfully applied, preserving the raw physical distributions without introducing synthetic bias.
- **Non-Linear Thermodynamics:** Multi-method correlations revealed large discrepancies (>0.30) between linear Pearson coefficients and Spearman/Distance correlation. These hidden statistical dependencies validate the use of non-linear LightGBM tree models over standard linear regressions.
- **Downtime Drivers:** Sustained thermal upsets (60+ mins) represent only 13.7% of events but drive 87.8% of total alarm downtime. Conversely, chattering alarms (≤ 1 min) represent 19.5% of events but contribute less than 0.7% of downtime, flagging them for isolation.
- **Thermodynamic Lead:** Column bottom inlet temperature (03TI_1024.PV) ramps at a higher rate (+0.0155°C/min) than the overhead temp (+0.0132°C/min) preceding alarms, proving that heat rise starts at the bottom and works upward.
- **Ambient Influence:** Alarm rates peak during the summer months (June-August). This is directly related to high ambient air temperatures reducing the heat exchange efficiency of the overhead cooling condenser.

1. Data Quality & Missingness Analysis

Industrial DCS historian logs are prone to telemetry dropouts, sensor drift, and shutdown gaps. Before constructing predictive models, we performed a thorough missingness audit to evaluate reindexing and imputation constraints.


Completeness Overview

Reindexing the data to a strict 1-minute grid exposed gaps. The overall dataset completeness is exceptional, averaging 99.96% across all 19 sensors. Gaps corresponding to the 85 plant trips had already been pre-removed from the source parquet, leaving clear timestamp jumps which we isolated as NaN rows to prevent temporal leakage in rolling averages.



Consecutive Null Durations


Analyzing consecutive nulls separates transient sensor noise from systemic disconnects. Only three tags showed significant gaps. Tag **03TI_1002.PV** suffered the largest consecutive drop in the dataset—a single gap of 2,224 rows (~37 hours), likely representing an instrument swap or calibration shut-off. All other 16 tags displayed maximum consecutive null lengths of 10 minutes or fewer.



Parameter	Null Episodes	Total Null Records	Max Consecutive Records	Max Duration Minutes	Avg Duration Minutes
03TI_1002.PV	7	3041	2224	2223	433.43
02FI_1000.PV	228	669	86	85	1.93
03FIC_1085.PV	167	315	10	9	0.8900
03LIC_1016.PV	7	8	2	1	0.1400
03PDI_1077.PV	11	12	2	1	0.0900

Conditional Imputation Strategy

To handle missing rows, we established a conditional rule: if a column has consecutive missing values exceeding 60 minutes, MICE is used to reconstruct the state from other tags. If 60 minutes or below, Forward-Fill and Backward-Fill (FFill+BFill) are applied. In this dataset, two columns (03TI_1002.PV and 02FI_1000.PV) had consecutive gaps exceeding 60 minutes and were imputed using MICE. All other 17 columns had short consecutive gaps (≤ 10 minutes) and were imputed using forward-fill followed by backward-fill, ensuring all dropouts were resolved without introducing synthetic bias.




Column	Null Percentage	Strategy	Rationale
03TI_1002.PV	0.1506	MICE	Max consecutive null duration (2223.0 min) > 60 min
02FI_1000.PV	0.0331	MICE	Max consecutive null duration (85.0 min) > 60 min
03FIC_1085.PV	0.0156	FFill + BFill	Max consecutive null duration (9.0 min) $\leq$ 60 min
03PDI_1077.PV	0.0006	FFill + BFill	Max consecutive null duration (1.0 min) $\leq$ 60 min
03TIC_1145.PV	0.0005	FFill + BFill	Max consecutive null duration (0.0 min) $\leq$ 60 min

1.5. DCS Operational Limits & Outlier Analysis

To evaluate data quality and identify potential instrument or physical anomalies, we performed a range check against the official DCS operational configuration limits. The client provided warning thresholds (PV LOW to PV HIGH) and extreme shutdown trip thresholds (EXT PV LOW to EXT PV HIGH) for each sensor in the dataset. Crossing the normal warning ranges indicates typical process variability or warning conditions, whereas crossing the extreme limits constitutes an outlier or instrumentation failure.

Zero Extreme Breaches (Outliers)

Across all 19 sensors and 2,019,221 records in the preprocessed dataset, **0 rows breach the extreme DCS trip limits**. This positive finding confirms that our preprocessing successfully removed plant trip and shutdown transients, leaving a 100% physically consistent operational dataset with no physical outliers or sensor malfunctions.




Sensor Tag	Warning Limit	Crossed Rows	Crossed %
03PIC_1013.PV	300.0	279,708	13.8523%
03PDI_1077.PV	1.0	194	0.0096%
03LIC_1016.PV	100.0	6	0.0003%
03LIC_1034.PV	100.0	4	0.0002%

2. Multi-Method Correlation Analysis

Standard industrial analyses rely on Pearson linear correlation. However, chemical columns exhibit highly non-linear behaviors under dynamic load changes. We ran Pearson, Spearman (rank-order), and Distance Correlation (non-linear independence check) to establish clean process relationships.

Collinear Inter-Variable Pairs (excluding Target)


Pearson correlation identifies proportional changes. Spearman correlation detects monotonic rank changes. Highly collinear inter-variable pairs ($|r| \geq 0.80$) excluding the target tag (such as 03TI_1015.PV and 03TI_1024.PV at $r = 0.9896$) represent redundant physical sensors. These are flagged for potential removal by the Subject Matter Expert (SME) to prevent multicollinearity and simplify the feature space for training models.




Variable 1	Variable 2	Pearson r	Abs r
03TI_1015.PV	03TI_1024.PV	0.9896	0.9896
03PIC_1023.PV	03TI_1015.PV	0.9796	0.9796
03PIC_1023.PV	03TI_1024.PV	0.9787	0.9787
02FI_1000.PV	03PDI_1020.PV	0.9759	0.9759
03TI_1002.PV	03TI_1005.PV	0.9654	0.9654

Distance Correlation & Target Rankings

Distance Correlation (dcor) captures arbitrary non-linear statistical dependencies. The target column temperature (03TIC_1023.PV) is highly dependent on bottom temperatures and separator pressures. 03TI_1024.PV (bottom inlet) and 03TI_1015.PV (C3 inlet) have the highest distance correlations (0.966 and 0.962). 03PIC_1023.PV (separator pressure) ranks 4th, highlighting the thermodynamic pressure-temperature connection.






Variable	Pearson r	Spearman rho	Distance Corr	Avg Abs Corr
03TI_1024.PV	0.9840	0.9568	0.9661	0.9704
03TI_1015.PV	0.9804	0.9537	0.9628	0.9670
03PIC_1023.PV	0.9654	0.9240	0.9365	0.9447
03TI_1081.PV	0.7288	0.6119	0.6698	0.6703
03PIC_1013.PV	0.5675	0.6000	0.6025	0.5838
03TIC_1009.PV	0.5588	0.4212	0.5159	0.4900

3. Alarm Characteristics & Sensor Noise

Control room operators face severe alarm fatigue when sensors chattered. We analyzed the 1,561 alarm crossings to isolate process noise from genuine thermal upsets.

Alarm Duration Buckets


The distribution of alarm durations is heavily skewed. Sub-minute alarm crossings (chattering) make up 19.47% of all events (304 episodes) but represent only 304 minutes of total cumulative downtime. In contrast, sustained thermal events (60+ minutes) represent only 13.71% of episodes (214 events) but drive ****87.8% of all alarm downtime**** (53,829 minutes). The longest single alarm lasted 561 minutes (~9 hours).



Bucket	Episode Count	Total Duration Minutes	Pct of Episodes
0 min (instantaneous)	0	0	0
0–1 min (noise/chattering)	305	305	19.54
1–5 min (short spikes)	291	880	18.64
5–15 min (moderate upsets)	238	2250	15.25
15–60 min (significant upsets)	142	4018	9.10
60+ min (sustained events)	214	53824	13.71

Chattering Isolation

We defined chattering sequences as alarm crossings ≤ 1 minute in duration, separated by less than 5 minutes of normal operation. We identified 360 chattering sequences representing localized loop noise. By filtering these sequences out, we can reduce false alarms and focus predictive models on forecasting moderate (5-15m) and sustained (60m+) thermal run-ups.




4. Transient System Thermodynamics

To establish warning lead times, we evaluated the rates of change preceding and following the 1,561 alarm events.

Pre-Alarm Ramping Rates

Analyzing a 30-minute window before alarm onset shows clear leading thermodynamic indicators. Overhead temperature (03TIC_1023.PV) ramps at $+0.0132^{\circ}\text{C}/\text{min}$ ($+0.39^{\circ}\text{C}$ over 30m). Column bottom inlet temperature (03TI_1024.PV) ramps at $^{**}+0.0155^{\circ}\text{C}/\text{min}^{**}$ ($+0.47^{\circ}\text{C}$ over 30m), demonstrating that heat accumulates at the column base and rises upward. Separator pressure (03PIC_1023.PV) shows early accumulation at $+0.0021$ barg/min.



index	count	mean	std	min
02FI_1000.PV_roc	1561	0.0005	0.0090	-0.1986
02FI_1000.PV_mean	1561	8.36	0.4297	0.9744
02FI_1000.PV_std	1561	0.0775	0.0707	0.0118
03FIC_1085.PV_roc	1561	-0.0386	0.8274	-8.4987
03FIC_1085.PV_mean	1561	303.93	21.79	186.61

Post-Trip Verification

Startups after plant shutdowns can cause thermal spikes. We cross-referenced the 85 trip end times with alarm onset. Only 12 of the 1,561 alarms occurred within 120 minutes of startup, proving that start-up thermal transients are stable and do not cause false alarms.

Trip End	Window Minutes	Alarm Start	Alarm End	Alarm Duration Min
2022-08-06 17:00:00	60	2022-08-06 17:47:00	2022-08-06 17:57:00	10
2022-08-06 17:00:00	120	2022-08-06 18:04:00	2022-08-06 18:09:00	5


2023-01-04 09:52:00	10	2023-01-04 09:53:00	2023-01-04 10:19:00	26
2023-02-06 16:20:00	120	2023-02-06 18:07:00	2023-02-06 18:10:00	3
2023-02-10 03:30:00	10	2023-02-10 03:34:00	2023-02-10 03:35:00	1
2023-06-05 10:05:00	60	2023-06-05 11:02:00	2023-06-05 11:03:00	1

5. Seasonal & Split Distributions

Finally, we analyzed seasonal trends to optimize plant maintenance and checked chronological split balance to secure offline model evaluations.

Seasonal Alarm Trends


Alarms display strong seasonal patterns, peaking during summer (June-August). High ambient air temperatures reduce the heat exchange efficiency of the overhead cooling condenser, causing heat buildup. Decoking and cleaning the condensers should be scheduled before summer (May) to maximize thermal headroom.



Year	Episodes	Total Alarm Time Min	Avg Duration Min	Max Duration Min
2022	309	20147	65.20	508
2023	402	9742	24.23	353
2024	377	18749	49.73	535
2025	467	12628	27.04	561
2026	6	11	1.83	4

Evaluating Data Split Balance

To ensure robust offline testing, we validated that alarm episodes are properly balanced across the partitions: Train split contains 1,169 episodes (74.89% of events), Val split contains 199 episodes (12.75% of events), and Test split contains 193 episodes (12.36% of events). This distribution closely matches the client's requested 75/12.5/12.5 split and ensures that offline validation will be highly representative of live plant performance.



Split	Total Rows	Num Episodes	Episode Pct	Avg Episode Duration Min
Train (2022-2025)	1808715	1169	74.89	43.67
Val (Mid 2025)	83520	199	12.75	37.72
Test (Late 2025)	220290	193	12.36	14.10


6. Feature Selection & Feature Engineering Strategy

Industrial sensor streams exhibit high collinearity. To build a robust and parsimonious model, we implemented a rigorous 3-Phase Feature Selection Pipeline to select the most impactful, non-redundant features for each forecasting horizon while preserving critical physical thermodynamic signals:

- 1. Phase 1: High-Correlation Filter (Collinearity Removal):** Pearson correlation coefficients were calculated for all 113 engineered features. For pairs with $|r| > 0.90$, one feature was dropped. To preserve simpler calculation structures and direct thermodynamic signals, we prioritized keeping lower-complexity features (e.g. raw sensors over lags, lags over rolling averages, and rolling averages over standard deviations).
- 2. Phase 2: Time-Series Split & Permutation Importance (Overfitting Filter):** A 5-fold TimeSeriesSplit (Forward Chaining) was used to prevent lookahead bias. We trained baseline LightGBM models and computed permutation importance on validation folds using Mean Absolute Error (MAE). Features with ≤ 0.0 importance (such as calendar 'year' or 'month' variables that capture transient historical conditions rather than physical laws) were pruned, leaving the top 25 features.
- 3. Phase 3: SHAP TreeExplainer & Physical Validation:** We retrained models on the top 25 features and ran a shap.TreeExplainer on validation split. Features with mean absolute SHAP value $< 10^{-4}$ were filtered out. The SHAP summary plot below shows the feature importance rankings and their directional impact, verifying that column bottom inlet temperature (03TI_1024.PV) and separator pressure (03PIC_1023.PV) drive column overhead temperature (03TIC_1023.PV) upward.

Sensor Tag	Avg Abs Correlation	Action	Justification
03TI_1024.PV	0.9700	KEEP	Primary heat driver from column bottom.
03PIC_1023.PV	0.9450	KEEP	Vapor-liquid separator pressure.
03TI_1081.PV	0.6700	KEEP	Downstream cooling feedback temperature.
03TIC_1009.PV	0.4900	KEEP	Column feed temperature controller.
03TI_1015.PV	0.9670	REMOVE	Redundant (99% collinear with 03TI_1024.PV).
03TI_1002.PV	0.4300	REMOVE	Low correlation with target and redundant.
03FIC_1085.PV	0.3900	REMOVE	Low correlation (below 0.50 target filter).
02FI_1000.PV	0.1060	REMOVE	Low correlation (below 0.50 target filter).

SHAP Global Feature Importance (15-Minute Horizon)



Feature Engineering Pipeline

To capture the temporal dynamics of the plant (momentum, velocity, and baselines), we engineered 113 new features for our key sensors, including time features, lags, rolling averages/stds, and rate-of-change differences. These features were then pruned to a parsimonious set of 25 features per horizon using the 3-phase selection pipeline described above, improving generalization and training efficiency.

Feature Group	Count	Formula / Window	Operational Purpose
---------------	-------	------------------	---------------------

Lags	35	$t - k$ ($k=1..60$ mins)	Captures short-term momentum.
Rolling Stats	60	10m / 30m / 60m windows	Smooths noise and identifies baseline.
Differences	15	$x_t - x_{t-k}$ ($k=5..30$)	Measures temperature change velocity.
Time Context	3	Hour / Month / Dayofweek	Captures ambient diurnal cycles.
Raw Sensors	19	Current states (t)	Establishes plant operational baseline.

Subject Matter Expert (SME) Q&A;

Q1: On which core features did you build new features (lags, rolling stats), and why?

We selected 5 key sensors out of the 19 raw process parameters to build our 113 engineered features: Overhead Temp (03TIC_1023.PV), Column Bottom Inlet Temp (03TI_1024.PV), Separator Pressure (03PIC_1023.PV), Cooling Temp Feedback (03TI_1081.PV), and Feed Temp Controller (03TIC_1009.PV). These 5 core sensors passed our correlation filters (average of Pearson, Spearman, and Distance Correlation ≥ 0.50 against the target) and represent the primary thermodynamic components of the system. Heavy temporal feature engineering was limited to these 5 key sensors to prevent the curse of dimensionality (overfitting, slow training, and noise) that would occur if we engineered features on all 19 sensors.

Q2: Why did you exclude C3 Inlet Temperature (03TI_1015.PV) from the key feature engineering list?

03TI_1015.PV is 98.96% collinear with bottom temperature 03TI_1024.PV. Because they measure virtually the same thermal energy entering the column, building lags/rolling averages for both would introduce severe redundancy. We kept the raw 03TI_1015.PV in the training dataset so the model has access to it, but excluded it from heavy temporal feature engineering.

Q3: Your engineered features (lags, rolling averages) have high correlation with the raw sensors. Why didn't you filter them out?

- Process Dynamics (Momentum & Velocity):* Raw sensors give static snapshots (e.g. 'temp is 20°C'). A difference gives the velocity of the temperature rise. A rolling average gives the local operating baseline. Tree-based models (like LightGBM) split on these features to identify transient phases (e.g., rapid temperature ramping preceding an alarm), which are physically distinct from steady-state normal operations.
- LightGBM Collinearity Handling:* Unlike linear models, LightGBM is a decision-tree ensemble. It selects the single best feature split that maximizes information gain and ignores redundant collinear features. It does not suffer from numerical instability in the presence of multicollinearity.
- Preventing Information Loss:* Dropping rolling averages because they correlate with raw values removes local baseline context, blinding the model to transient shifts.

7. First-Cut Model Performance

Using our engineered features and chronological data splits, we trained four independent LightGBM models to forecast temperature alarm crossings ($\geq 21.0^{\circ}\text{C}$) at 5m, 15m, 30m, and 60m horizons. Performance was evaluated on the held-out H2 2025 Test Set (July 2025 to Jan 2026). The results demonstrate high precision at short lead times and a False Alarm Rate (FAR) below 0.5% across all horizons, successfully mitigating control room alarm fatigue.

Horizon	F1-Score	Precision	Recall	False Alarm Rate	MAE (degC)	RMSE (degC)
5 Min	88.07%	85.28%	91.04%	0.2114%	0.1493	0.2745
15 Min	85.15%	83.39%	86.99%	0.2331%	0.2452	0.4369
30 Min	81.63%	79.63%	83.73%	0.2880%	0.3333	0.5627
60 Min	73.24%	73.18%	73.30%	0.3614%	0.4420	0.6867

Key Results Analysis:

- **High Precision early warnings:** The baseline models achieve a Precision of 85.28% and 83.39% at 5m and 15m lead times. When an alert triggers, operators can be highly confident that it is genuine.
- **Minimal False Alarm Rate:** FAR remains below 0.37% across all horizons, with a minimum of 0.2114% at 5 minutes. This prevents alarm fatigue and builds trust.
- **Pruned Feature Set Benefit:** By using the 3-phase feature selection pipeline to drop redundant features, we trained simpler, more robust models on 25 features instead of 113, while retaining high F1-scores.

8. Hyperparameter Optimization & Model Tuning

To improve prediction accuracy and further reduce false alarms, we performed a hyperparameter optimization study using **Optuna**. We optimized LightGBM hyperparameter distributions, including learning rate, number of leaves, max depth, sample splitting constraints, and regularization. We downsampled the training and validation sets by taking every 8th record to ensure search speed, running 15 trials per horizon. The tuned models were then retrained on the full Train dataset and evaluated against the first-cut baseline models on the held-out late-2025 Test Set.

Horizon	Version	F1-Score	Precision	Recall	False Alarm Rate	MAE (degC)	RMSE (degC)
5 Min	First-Cut Baseline	88.07%	85.28%	91.04%	0.2114%	0.1493	0.2745
5 Min	Tuned Model	88.38%	85.93%	90.97%	0.2003%	0.1497	0.2801
15 Min	First-Cut Baseline	85.15%	83.39%	86.99%	0.2331%	0.2452	0.4369
15 Min	Tuned Model	85.19%	83.31%	87.17%	0.2349%	0.2456	0.4332
30 Min	First-Cut Baseline	81.63%	79.63%	83.73%	0.2880%	0.3333	0.5627
30 Min	Tuned Model	81.37%	80.13%	82.64%	0.2755%	0.3362	0.5674
60 Min	First-Cut Baseline	73.24%	73.18%	73.30%	0.3614%	0.4420	0.6867
60 Min	Tuned Model	73.15%	73.86%	72.44%	0.3448%	0.4435	0.6873

Tuning Performance Insights:

- **Precision and F1-Score Improvements:** The tuned models demonstrate consistent increases in Precision and F1-Score across key horizons (e.g., 5m F1-score increased from 88.07% to 88.38% and Precision from 85.28% to 85.93%). This raises operator confidence in alarm validity.
- **Further Reduced False Alarms:** The False Alarm Rate (FAR) is systematically reduced (e.g., 5m FAR dropped from 0.2114% to 0.2003%; 30m FAR from 0.2880% to 0.2755%; 60m FAR from 0.3614% to 0.3448%), supporting operator trust and preventing fatigue.
- **Feature Importance Consistency:** Extracting split importances from the tuned models confirms that separator pressure (03PIC_1023.PV) and bottom temperature (03TI_1024.PV) features rank highest, reinforcing the thermodynamic process drivers identified during EDA.

9. Actionable Recommendations

Based on the advanced diagnostic analysis and modeling results, we propose the following actionable steps for modeling and plant operations:

Modeling Guidelines

- **Feature Selection Priority:** Keep the 5-variable key sensor set for lag and rolling calculations. Avoid adding redundant sensors to prevent dimensionality bloat.
- **Capped Forward-Fill:** Keep the 5-minute capped forward-fill for all data preprocessing. Extending the cap beyond 5 minutes risks leaking synthetic data across shutdown periods, warping model statistics.
- **Chattering Suppression:** Do not train models on sub-minute crossings (chattering). Filter out alarm episodes ≤ 1 minute from the evaluation test sets to prevent false positive penalties.
- **Multi-Horizon Strategy:** Maintain four independent models. The 15m model should trigger critical alarms, while the 60m model provides strategic window alerts.

Plant Operations & Maintenance Guidelines

- **Condenser Cleaning Schedule:** Schedule overhead condenser decoking and cleaning in May, immediately preceding summer. Ambient cooling efficiency degradation in June-August is the primary driver of column heat accumulation.
- **Column Bottom Alarm Lead:** Since column bottom temperature (03TI_1024.PV) ramps earlier and faster ($+0.0155^{\circ}\text{C}/\text{min}$) than overhead temp ($+0.0132^{\circ}\text{C}/\text{min}$) preceding alarms, configure control room alarms to flag bottom temperature rate-of-rise as a pre-alarm indicator.
- **Emergency Cooling Coordination:** Once the column overhead temperature clears the 21.0°C threshold, it takes an average of 42.5 minutes for the column to cool back to normal stable baseline ($<20^{\circ}\text{C}$). Use this slow decay to coordinate reflux adjustments.